

Acceptance Deck

River Hawk Consulting, LLC · IMAX HCD Prototype OT
August 7, 2026

CONTENTS

IMAX Chat-Viewer Prototype OT — Acceptance Deck (DRAFT for internal feedback)

- 1. Objectives (from the solicitation)
- 2. What was built
- 3. Test results (evidence: project/evidence/)
- 4. Open issues (contractually honest — deliverable ⑥ expects these)
- 4a. How to watch the demo
- 5. Ask of reviewers (this draft)

IMAX Chat-Viewer Prototype OT — Acceptance Deck (DRAFT for internal feedback)

River Hawk Consulting, LLC · UNCLASSIFIED // Proprietary Period of performance: **6 – 20 August 2026** (executed 6 August) · SPA build complete 1 August, ahead of award (`v1.0.0-prototype`) Current package: `v1.1.0-prototype` — adds the Platforma environment work and completes Deliverable ① *For review and direction — not a delivery transmittal. The prototype has not run in Platforma; sandbox access has not been granted.*

1. Objectives (from the solicitation)

Deliverable	Objective	Status
① HCD artifacts	Personas, empathy maps, wireframes grounding the design	✅ complete + three workflow-model findings
② Angular SPA codebase	The evaluated chat-viewer, rebuilt in Angular, in GitHub	✅ feature-complete on <code>main</code>
③ e2e smoke tests	The Phase 2 acceptance suite against the new build	✅ 7/7 flows covered (3 verbatim, 4 amended) + 25 more
④ User + developer guides	Workflow-organized and architecture docs	✅ complete, PDFs in <code>project/delivery/</code>
⑤ Sprint artifacts	Board, burndown, risk register — worked live	✅ maintained daily, final exports at freeze
⑥ Demo & acceptance package	Recorded walkthrough, this deck, exports	🟡 this draft

2. What was built

- **Angular 21 (LTS, pinned)** SPA at `angular/` — standalone components, signals, typed contracts against the frozen mock services. The React demo remains in-repo as the reference implementation; `data-testid`s carry 1:1, which is the acceptance mechanism.
- **Full evaluated workflow:** triage queue with My-languages scope and thread-level dispositions; corpus-wide 4-language search (content + entity-selector modes, facets, geo-fence and watchlist groups, date range); jump-to-evidence with flash; translate-in-place with linguist review verdicts; entity chips; analyst notes; OCR viewer with officer annotation and the tag-facet loop-back; gold-copy tray with provenance and standardized export.
- **Three HCD findings (Sponsor validation pending — see the demo's before/after segments):**
 1. *The unit of gold is the whole thread* — translate-thread, gold-ready meter, promote-to-gold (full verdicted transcript), and the queue's N-of-M/stack-clear affordances (`hcd/linguist_workflow_model.md`).
 2. *No verdict without the source* — always bilingual (English leads, original stays on screen, direction-correct); corrections are written against the source (`hcd/bilingual_display_model.md`).
 3. *A linguist's bench, one output* — one gold action, thread-level enrichment, a button-free stack, collapsible find tools (`hcd/one_output_model.md`). **This one removes evaluated behavior** (snippet clipping); the React reference build is retained and the removal reverts in a single commit.
- **Platforma packaging:** six images (unprivileged nginx SPA + five single-responsibility service pods), kustomize manifests with probes/limits/non-root, no secrets, stateless. Proven on a kind cluster in CI on every push (build → trivy CRITICAL gate → deploy → rollout → in-cluster smoke).

3. Test results (evidence: `project/evidence/`)

Suite	Result
Ported Phase 2 acceptance flows	7 of 7 covered — 3 verbatim, 4 amended by logged HCD decisions
Total e2e (Playwright, vs live mocks)	32 green (JSON results in evidence)
Angular unit tests (wire-format + shell)	10 green
CI	lint · prod build · unit · e2e · package-validate — green on <code>main</code>
Security gate	trivy caught 2 real CRITICALs during hardening; both fixed; gate green
Requirements compliance	Every solicitation requirement scored — <code>requirements_compliance.md</code>

4. Open issues (contractually honest — deliverable ⑥ expects these)

1. **Sponsor unresponsive since kickoff** — proceeding on stated assumptions (Angular v21 LTS; deliver-deployable). Escalation: kickoff email sent 31 Jul + phone follow-up.
2. **Thread-gold amendment awaits Sponsor validation** (touchpoint #1, Mon 3 Aug). Removable without touching evaluated behavior if strict parity is preferred.
3. **Platforma sandbox access not granted** — validation is on kind in CI (the agreed fallback); sandbox deployment is a same-day exercise once credentials arrive.
4. **Demo narration is complete** — the walkthrough was re-captured on 7 August against the fixed build and narrated (4:08). The synthesized scratch track is superseded. Script, narration text, capture pipeline and build are all in the repository, so the recording is reproducible rather than a one-off.
5. **Live model gateway is wired but unproven against the enclave** — configuration-only switch (`deploy/README.md` §4); cannot be exercised without access.

6. **Screenshot-diff specs** intentionally trimmed per the pre-agreed scope-trim order; parity is evidenced by side-by-side captures instead.
7. **1 Aug design decisions (owner-directed, logged in the HCD artifacts)** — the defining one: **this application is a linguist's bench** (`hcd/one_output_model.md`). Translations render by default (bilingual; frontier-LLM retranslation is the named production seam); review state as colored message edges; ONE output — promoted thread gold (snippet evidence clips removed; analyst notes travel into the gold transcript); enrichment is thread-level, run once (translate + entity extraction); queue offers only *reviewed* — flag/discard live in the viewer strip; gold pane slimmed. Acceptance posture: specs 1/3/4/5 verbatim, 2/6/7 amended with the beats retained where behavior survives; React reference app untouched; single-commit restoration path if the Sponsor requires strict parity. A multi-page Russian document (3-page storage contract) was added to the corpus with OCR page navigation.

4a. How to watch the demo

`RiverHawk_IMAX_Design_Evolution_and_Demo.mp4` (1:57) is built for this review. The first 69 seconds show each HCD finding as a **live before/after** — the July build you evaluated runs, then the delivered build runs on the same thread — so the changes can be judged on their merits rather than discovered. The remainder is the working product end to end. A capability-only cut is included separately for audiences who don't need the derivation.

5. Ask of reviewers (this draft)

- Does the thread-gold flow read as *the* linguist workflow, or does anything about the meter/promote/stack semantics feel off? (This is the HCD story for touchpoint #1.)
- Bilingual view: is English-leads-source-beneath the right hierarchy, or would a side-by-side split serve the linguist better? (Finding #2 — shown as before/after in the demo and in `../evidence/screens/day1-bilingual-review.png`.)
- Finding #3 removes snippet clipping. Is the reasoning persuasive, or should the evaluated behavior be restored?
- Any gaps between the demo and how you'd present this to the Sponsor?
- Deck/guide tone: anything that reads as overclaiming or underclaiming?