

# Acceptance Deck

---

**River Hawk Consulting, LLC · IMAX HCD Prototype OT**  
August 8, 2026

# CONTENTS

## IMAX Chat-Viewer Prototype OT — Acceptance Deck (DRAFT for internal feedback)

- 1. Objectives (from the solicitation) .....
- 2. What was built .....
- 3. Test results (evidence: project/evidence/) .....
- 4. Open issues (contractually honest — deliverable ⑥ expects these) .....
- 4a. How to watch the demo .....
- 5. Ask of reviewers (this draft) .....

# IMAX Chat-Viewer Prototype OT — Acceptance Deck (DRAFT for internal feedback)

**River Hawk Consulting, LLC** · UNCLASSIFIED // Proprietary Period of performance: **6 – 20 August 2026** (executed 6 August) · SPA build complete 1 August, ahead of award ( `v1.0.0-prototype` ) Current package: `v1.1.0-prototype` — adds the Platforma environment work and completes Deliverable ① *For review and direction — not a delivery transmittal. The prototype has not run in Platforma; sandbox access has not been granted.*

## 1. Objectives (from the solicitation)

| Deliverable                 | Objective                                                | Status                                                |
|-----------------------------|----------------------------------------------------------|-------------------------------------------------------|
| ① HCD artifacts             | Personas, empathy maps, wireframes grounding the design  | ✅ complete + <b>three</b> workflow-model findings     |
| ② Angular SPA codebase      | The evaluated chat-viewer, rebuilt in Angular, in GitHub | ✅ feature-complete on <code>main</code>               |
| ③ e2e smoke tests           | The Phase 2 acceptance suite against the new build       | ✅ 7/7 flows covered (3 verbatim, 4 amended) + 25 more |
| ④ User + developer guides   | Workflow-organized and architecture docs                 | ✅ complete, PDFs in <code>project/delivery/</code>    |
| ⑤ Sprint artifacts          | Board, burndown, risk register — worked live             | ✅ maintained daily, final exports at freeze           |
| ⑥ Demo & acceptance package | Recorded walkthrough, this deck, exports                 | 🟡 this draft                                          |

## 2. What was built

- **Angular 21 (LTS, pinned)** SPA at `angular/` — standalone components, signals, typed contracts against the frozen mock services. The React demo remains in-repo as the reference implementation; `data-testid`s carry 1:1, which is the acceptance mechanism.
- **Full evaluated workflow:** triage queue with My-languages scope and thread-level dispositions; corpus-wide 4-language search (content + entity-selector modes, facets, geo-fence and watchlist groups, date range); jump-to-evidence with flash; translate-in-place with linguist review verdicts; entity chips; analyst notes; OCR viewer with officer annotation and the tag-facet loop-back; gold-copy tray with provenance and standardized export.
- **Three HCD findings (Sponsor validation pending — see the demo's before/after segments):**
  1. *The unit of gold is the whole thread* — translate-thread, gold-ready meter, promote-to-gold (full verdicted transcript), and the queue's N-of-M/stack-clear affordances ( `hcd/linguist_workflow_model.md` ).
  2. *No verdict without the source* — always bilingual (English leads, original stays on screen, direction-correct); corrections are written against the source ( `hcd/bilingual_display_model.md` ).
  3. *A linguist's bench, one output* — one gold action, thread-level enrichment, a button-free stack, collapsible find tools ( `hcd/one_output_model.md` ). **This one removes evaluated behavior** (snippet clipping); the React reference build is retained and the removal reverts in a single commit.
- **Platforma packaging:** six images (unprivileged nginx SPA + five single-responsibility service pods), kustomize manifests with probes/limits/non-root, no secrets, stateless. Proven on a kind cluster in CI on every push (build → trivy CRITICAL gate → deploy → rollout → in-cluster smoke).

### 3. Test results (evidence: `project/evidence/`)

| Suite                                    | Result                                                                          |
|------------------------------------------|---------------------------------------------------------------------------------|
| Ported Phase 2 acceptance flows          | <b>7 of 7 covered</b> — 3 verbatim, 4 amended by logged HCD decisions           |
| Total e2e (Playwright, vs live mocks)    | <b>32 green</b> (JSON results in evidence)                                      |
| Angular unit tests (wire-format + shell) | 10 green                                                                        |
| CI                                       | lint · prod build · unit · e2e · package-validate — green on <code>main</code>  |
| Security gate                            | trivy caught 2 real CRITICALs during hardening; both fixed; gate green          |
| Requirements compliance                  | Every solicitation requirement scored — <code>requirements_compliance.md</code> |

### 4. Open issues (contractually honest — deliverable ⑥ expects these)

1. **Sponsor unresponsive since kickoff** — proceeding on stated assumptions (Angular v21 LTS; deliver-deployable). Escalation: kickoff email sent 31 Jul + phone follow-up.
2. **Thread-gold amendment awaits Sponsor validation** (touchpoint #1, Mon 3 Aug). Removable without touching evaluated behavior if strict parity is preferred.
3. **Platforma sandbox access not granted** — validation is on kind in CI (the agreed fallback); sandbox deployment is a same-day exercise once credentials arrive.
4. **Demo narration is complete** — the walkthrough was re-captured on 7 August against the fixed build and narrated (4:08). The synthesized scratch track is superseded. Script, narration text, capture pipeline and build are all in the repository, so the recording is reproducible rather than a one-off.
5. **Live model gateway is wired but unproven against the enclave** — configuration-only switch (`deploy/README.md` §4); cannot be exercised without access.

6. **Screenshot-diff specs** intentionally trimmed per the pre-agreed scope-trim order; parity is evidenced by side-by-side captures instead.
7. **1 Aug design decisions (owner-directed, logged in the HCD artifacts)** — the defining one: **this application is a linguist's bench** (`hcd/one_output_model.md`). Translations render by default (bilingual; frontier-LLM retranslation is the named production seam); review state as colored message edges; ONE output — promoted thread gold (snippet evidence clips removed; analyst notes travel into the gold transcript); enrichment is thread-level, run once (translate + entity extraction); queue offers only *reviewed* — flag/discard live in the viewer strip; gold pane slimmed. Acceptance posture: specs 1/3/4/5 verbatim, 2/6/7 amended with the beats retained where behavior survives; React reference app untouched; single-commit restoration path if the Sponsor requires strict parity. A multi-page Russian document (3-page storage contract) was added to the corpus with OCR page navigation.

## 4a. How to watch the demo

Two cuts ship, and only two:

- `RiverHawk_IMAX_Walkthrough_2026-08-08.mp4` (4:05, narrated) — the working product end to end, with each design finding explained where it shows on screen: why the unit of gold is the whole thread, why translations are always bilingual, why the queue asks one question. Watch this one.
- `RiverHawk_IMAX_Walkthrough_2026-08-08_capability.mp4` (3:49) — the same footage without the design derivation, for audiences who want the capability and not the reasoning.

Both are reproducible rather than one-offs: `narration.json` /

`narration_capability.json` carry the script, `scripts/capture-walkthrough.js` drives the capture, and `scripts/build-walkthrough.sh` builds either cut.

## 5. Ask of reviewers (this draft)

- Does the thread-gold flow read as *the* linguist workflow, or does anything about the meter/promote/stack semantics feel off? (This is the HCD story for touchpoint #1.)

- Bilingual view: is English-leads-source-beneath the right hierarchy, or would a side-by-side split serve the linguist better? (Finding #2 — shown as before/after in the demo and in `../evidence/screens/day1-bilingual-review.png`.)
- Finding #3 removes snippet clipping. Is the reasoning persuasive, or should the evaluated behavior be restored?
- Any gaps between the demo and how you'd present this to the Sponsor?
- Deck/guide tone: anything that reads as overclaiming or underclaiming?