

Acceptance Deck

River Hawk Consulting, LLC · IMAX HCD Prototype OT
August 8, 2026

CONTENTS

IMAX Chat-Viewer Prototype OT — Acceptance Deck

1. Objectives (from the solicitation)

2. What was built

3. Test results (evidence: project/evidence/)

4. Open issues

5. How to watch the demo

IMAX Chat-Viewer Prototype OT — Acceptance Deck

River Hawk Consulting, LLC · UNCLASSIFIED // Proprietary Period of performance: **6 – 20 August 2026** (executed 6 August) · SPA build complete 1 August, ahead of award (v1.0.0-prototype) Current package: **v1.3.0-prototype** — the delivered acceptance package *Transmitted for acceptance at M5. The prototype has been deployed and run in Platforma; the ArgoCD capture evidencing it is attached to the high-side repository, which is where it is producible.*

1. Objectives (from the solicitation)

Deliverable	Objective	Status
① HCD artifacts	Personas, empathy maps, wireframes grounding the design	✅ complete + three workflow-model findings
② Angular SPA codebase	The evaluated chat-viewer, rebuilt in Angular, in GitHub	✅ feature-complete on <code>main</code>
③ e2e smoke tests	The Phase 2 acceptance suite against the new build	✅ 7/7 flows covered (3 verbatim, 4 amended); 72 e2e specs in total
④ User + developer guides	Workflow-organized and architecture docs	✅ complete, PDFs in <code>project/delivery/</code>
⑤ Sprint artifacts	Board, burndown, risk register — worked live	✅ maintained daily, final exports at freeze
⑥ Demo & acceptance package	Recorded walkthrough, this deck, exports	✅ complete

2. What was built

- **Angular 21 (LTS, pinned)** SPA at `angular/` — standalone components, signals, typed contracts against the frozen mock services. The React demo remains in-repo as the reference implementation; `data-testid`s carry 1:1, which is the acceptance mechanism.
- **Full evaluated workflow:** triage queue with My-languages scope and thread-level dispositions; corpus-wide 4-language search (content + entity-selector modes, facets, geo-fence and watchlist groups, date range); jump-to-evidence with flash; translate-in-place with linguist review verdicts; entity chips; analyst notes; OCR viewer with officer annotation and the tag-facet loop-back; gold-copy tray with provenance and standardized export.
- **Three HCD findings (Sponsor validation pending — see the demo's before/after segments):**
 1. *The unit of gold is the whole thread* — translate-thread, gold-ready meter, promote-to-gold (full verdicted transcript), and the queue's N-of-M/stack-clear affordances (`hcd/linguist_workflow_model.md`).
 2. *No verdict without the source* — always bilingual (English leads, original stays on screen, direction-correct); corrections are written against the source (`hcd/bilingual_display_model.md`).
 3. *A linguist's bench, one output* — one gold action, thread-level enrichment, a button-free stack, collapsible find tools (`hcd/one_output_model.md`). **This one removes evaluated behavior** (snippet clipping); the React reference build is retained and the removal reverts in a single commit.
- **Platforma packaging:** six images (unprivileged nginx SPA + five single-responsibility service pods), kustomize manifests with probes/limits/non-root, no secrets, stateless. Proven on a kind cluster in CI on every push (build → trivy CRITICAL gate → deploy → rollout → in-cluster smoke).

3. Test results (evidence: `project/evidence/`)

Suite	Result
Ported Phase 2 acceptance flows	7 of 7 covered — 3 verbatim, 4 amended by logged HCD decisions
Total e2e (Playwright, vs live mocks)	72 green (JSON results in evidence)
Angular unit tests (wire-format + shell)	10 green
Node unit tests (identity, certificates, OCR fixtures, model gateway, route smoke)	93 green
Interaction latency at volume	~17 ms worst scroll-to-repaint against the TDD's 200 ms budget, measured on the 2,388-message standing channel
CI	lint · prod build · unit · e2e · package-validate — green on <code>main</code>
Security gate	trivy caught 2 real CRITICALs during hardening; both fixed; gate green
Requirements compliance	Every solicitation requirement scored — <code>requirements_compliance.md</code>

4. Open issues

Stated rather than smoothed. Nothing here blocks acceptance of the delivered artifacts; each is either Government-side, environment-dependent, or a scope boundary that was decided deliberately and is recorded where the decision was made.

Government-side or environment-dependent

- **Deployment evidence (M2) lives high-side.** The application is delivered as a deployable image set with an ArgoCD Application manifest and a Helm chart / kustomize pair proven equivalent in CI, and it has been deployed and run in Platforma. The sync, health and resource-tree record is generated inside the enclave, so it is attached to the high-side repository rather than to this package — River Hawk has no low-side path to carry it.
- **A live model gateway is built beyond the requirement,** and has not been called from inside the enclave. The solicitation accepts containerized mock services that emulate production, and the five schema-pinned services satisfy it; the gateway path is capability added on top. Translation, entity extraction, summarization and OCR are fixture-backed by default and gateway-backed by configuration — `MODEL_ENDPOINT`, `MODEL_NAME` and `MODEL_API_KEY` only, no client rework. The client side is proven: the provider was run against a real OpenAI-compatible model and the wire-format differences a stub cannot surface were found and handled (`max_tokens` vs `max_completion_tokens`; reasoning models rejecting `temperature`), both negotiated at pod startup. What remains is pointing it at the in-enclave endpoint. Carried as risk 7.
- **Administrative deliverables.** Disposition of the FOCI package (Article XIV(c)) and the Cleared Personnel Certification and Disposition Reports (Article XV(n)/(o)) is for the Agreements Officer to confirm. All work is unclassified and no personnel are cleared under this agreement, so nil returns are expected — but a nil return is still a return.

Scope boundaries, decided and recorded

- **No media-query breakpoints.** The layout is fluid, the panels resize and collapse, and the scan viewer maximizes — but the build targets the analyst desktop it was evaluated on. Small-screen breakpoints were not in scope and are not claimed.
- **No client-side persistence.** Panel widths, sort order, verdicts, notes and tags live in memory and reset on reload, per TDD Task 3. `SessionStore` is the seam if production requires otherwise.
- **Snippet evidence clips were removed** on 1 August by HCD finding #3 (`hcd/one_output_model.md`), which found two competing output paths for one user. The reasoning and a single-commit restoration path are both recorded; the React reference build at `src/` retains the evaluated behaviour unchanged.

Delivered on the day

- **The live demonstration** is provided in person on 20 August. The narrated walkthrough is delivered alongside it as the standing record and as the fallback if enclave access is unavailable on the day.

5. How to watch the demo

One recording, which is the deliverable:

`RiverHawk_IMAX_Walkthrough_2026-08-08.mp4` (4:51, narrated) — the working product end to end, with each design finding explained where it shows on screen: why the unit of gold is the whole thread, why translations are always bilingual, why the queue asks one question.

It is reproducible rather than a one-off: `narration.json` carries the script, `scripts/capture-walkthrough.js` drives the capture, and `scripts/build-walkthrough.sh` assembles it.