

User Guide

River Hawk Consulting, LLC · IMAX HCD Prototype OT

August 8, 2026

CONTENTS

User Guide — IMAX Triage Workspace (the Linguist's Bench)	
The bench at a glance	
Working the stack — the core loop	
The workflow strip — every thread decision in one place	
Documents	
Finding things	
For supervisors	

User Guide — IMAX Triage Workspace (the Linguist's Bench)

UNCLASSIFIED — Deliverable ④. **This application is a linguist's bench**

(`../hcd/one_output_model.md`): you work a finite stack of threads and documents in your languages, judge every translation against its source, and produce one thing — the gold copy of the whole thread. Screens referenced here have live captures in `../project/evidence/screens/`. All data is fabricated demonstration content, marked by the banner on every screen.

The bench at a glance

Search & Triage (left): your stack — the queue, scoped to *My languages*, sorted oldest-first, with the `N/M worked` progress bar — plus corpus-wide search when you need to find something. The rows carry no buttons: a thread leaves your stack when you finish it, not when you tick a box. The magnifier in the panel header hides the search and filter tools entirely when you just want to work the queue — hiding them also clears any active search, so you always land back on your own stack. **Chat Log Viewer** (center): the reading and judging surface. **Gold Copy** (right, slim, collapsible): your finished output. One workflow strip above the thread carries every thread-level action.

Both side panels **resize**: drag the strip on the edge each one faces the centre across, or use the arrow keys once it has focus. Double-click it to go back to the default width. A wide manifest or a long thread can have the room it needs. Widths last for the session — this application stores nothing between reloads, deliberately, so a refresh puts everything back where it started.

Working the stack — the core loop

1. **Open a thread.** It arrives ready: every foreign message is already translated and shown **bilingually** — English leads, the original stays beneath it in its own script and direction — always. There is no show/hide toggle: you never judge blind, and you never click to see a translation.

2. **Judge each message.** ✓ **confirm** the machine translation or ✎ **correct** it — the editor shows the source directly above your draft. The message's edge shows its state at a glance: green = confirmed, rose = your correction, dashed = still waiting on you. The meter in the strip counts you toward **GOLD-READY**.

A verdict can be taken back. Click ✓ again on a message you have confirmed and it returns to the machine translation — a mis-click is not a decision you are stuck with. Clearing a *correction* discards what you wrote, so that one asks first.

3. **Add notes where the machine can't.** Code words, source correlation, context — the note displays under the message and travels into the gold transcript as a NOTE [analyst]: line.

A note on the conversation itself is a different thing, and has its own place. *Add conversation note* in the workflow strip records your read on the exchange as a whole — “this doesn't sound like a native Russian speaker”, “these two clearly know each other” — the judgement that comes from having read the thread through rather than from any single message. It sits above the stream and **leads the gold copy**, above the transcript, because it is a conclusion about all of it. Attached to a message it would be read as evidence from that message.

4. **Promote to gold.** When every foreign message carries your verdict, *Promote to gold* produces the thread's gold copy — the full transcript, original and English side by side with verdicts and notes — in the right pane, and marks the thread worked. Your stack ticks down; the queue tells you when it's clear.
5. **Export** renders your finished gold copies as full-translation products, each line traceable. **Copying confirms itself** — if the clipboard is unavailable the workspace says so and offers to select the text instead, rather than failing silently.

Reading order. The stream is sorted oldest-first and the header toggles it to newest-first. Order is decided here rather than taken from the feed, because collection does not guarantee chronological delivery. A promoted gold copy is always chronological whichever way you were reading — a record read backwards misleads whoever receives it.

Long threads. The stream renders only what is on screen, so a thread of any length opens and scrolls at the same speed.

Ordering your gold copies. Drag the handle on any promoted thread to reorder the tray. The export follows that order, so when a product draws on several threads the sequence is yours to decide.

The workflow strip — every thread decision in one place

- **Flag for targeter** — hand the thread to targeting without leaving your loop.
- **Discard** — a deliberate decision made *after reading*, from here, never from the queue (the stack is assigned work).
- **Extract entities** — one time, for the whole thread: names, places, numbers, and passports light up as colored chips on each message. Spotting aids, not buttons.
- **Translate *N* remaining** — the button counts the messages still missing a translation and re-runs the baseline batch for exactly those (e.g. late arrivals). A premium frontier-model retranslation is planned at the production cutover.
- **Accept remaining *N*** — on a long conversation only. A standing channel runs to thousands of messages, and clearing the tail one click at a time is not a workflow, so this accepts the outstanding machine translations in a single action. It tells you how many before it does anything, and **it does not claim you read them**: those messages are recorded in the gold copy as `bulk-accepted`, not `linguist-confirmed`, so whoever receives the transcript can tell the two apart. On a short thread the button is not offered at all — there, reading each message *is* the job.
- **Add conversation note** — your read on the exchange as a whole (see the core loop above).
- **Promote to gold** — the one gold action, locked until the meter reads GOLD-READY.

Documents

Threads with a paperclip carry an embedded file; document-lane items *are* files. The OCR viewer shows the scan above the extracted text — multi-page documents page both together — with a whole-document English gloss beneath.

Each line carries its own English, beside it. The extracted text is numbered line by line, and every line is paired with its translation on the same row, so you read across one line instead of hunting between two blocks. **The original sits on the side its script starts from** — left for Russian and Chinese, right for Arabic and Farsi — so a scan always reads outward from the margin it occupies on the page. A tabular form, where the manifest rows matter more than the prose, is where this earns its keep.

Make it big enough to read. The viewer maximizes to nearly the full screen, resizes from its bottom-right corner, and the divider between the scan and the text below it drags up and down, so you can give either one more of the window. Zoom offers **fit page width, fit whole page, actual size and 200%**, and a zoomed scan pans inside its pane. A dense form — the five-page Arabic customs declaration in the corpus is the example — is meant to be read at size, not squinted at.

The **officer annotation** strip is yours: quick tags (`priority`, `identity-doc`, `matches-open-case`, `follow-up`) or free-text tags, plus a document note. Type a new tag and press **Add** — or Enter, or simply click away; all three commit it. **A tag you coin is offered on every document afterwards**, lit when it applies to the one in front of you, because the whole point of a tag is gathering documents under it. Tags loop back into Search & Triage as the **Officer tags** facet and filter your queue to threads carrying that mark (a saved view — worked items stay visible under it).

Finding things

Type two characters and the search sweeps every message in the corpus, any language, *including the English translations* of foreign originals (`matched in EN translation` badge). Entity mode hunts detected entities — people, places, phone numbers and passports — and also matches sender handles, so a selector query finds the account as well as the number. Geo-fence and watchlist **groups** are one-click scopes; the facet row (Selector, Passport, Person, Geo, Image, Unenriched) narrows to messages carrying that signal; the date range composes with everything. **Flagged** appears with a count once you have handed anything to targeting, and isolates that pile — a handoff you cannot find again is a handoff you have lost. *Flagged first* in the sort list lifts them without hiding the rest. Click a hit and the thread opens scrolled to the evidence, flashed so you don't lose it. **Summarize** gives you an on-demand thread gist that collapses back out of the way.

For supervisors

The queue header answers the morning question: `Queue · N new · M worked`, the progress bar, and content-type lanes with corpus counts. Sorting offers oldest-first, most recent, message count and flagged-first. (Untranslated-first was retired on 1 August when threads began arriving already translated — there is no untranslated pile left to lift.) Review edges make any thread's progress readable over a shoulder. Downstream roles consume the bench's outputs — flagged threads (targeting) and exported gold copies (production) — rather than working inside it.